



Technical Skills

- **Code:** Python, bash, SQL, Markdown, powershell, Java, C/C++
- **Config:** cloud-init, SaltStack
- **CI/CD:** Github Actions, Azure DevOps, Jenkins, ArgoCD
- **Container:** Docker, Podman, Kubernetes(kOps, kubectl, helm)
- **Cloud/IaC:** AWS, Azure, Terraform, Packer, Ansible
- **Observability:** OpenTelemetry, CloudWatch, Datadog, Grafana

Experience

Software Engineer(DevOps)

July 2021 - Present

DataJoint - Science Operation for Neuroscience. [[Open-source](#)] and [[Commercial](#)]

Houston, TX

- * **AWS:** Administrated DataJoint's AWS account and several other customers' AWS accounts. Configured **VPC**, **Subnet**, **Security Groups**, **IAM** role and policies, **S3** lifecycle management, **EFS** access point, **EC2** instances, **RDS** instances, **Lambda** triggered by **SQS** or **EventBridge**, **SNS** and **SES**, **CloudWatch** metrics and alarms, **Route 53** DNS records, **Secrets Manager** for deployment secrets.
- * **CI/CD:** Developed generic **Github Actions** reusable workflows used by **30+** repositories followed by [Conventional Commits](#), [Release Flow](#) and [GitOps](#) best practices, to automate build, test, release, publish private or open-source **Python** packages[[PyPI](#)] or deploy **Docker** images[[Dockerhub](#)].
- * **Kubernetes:** Provisioned Kubernetes clusters hosted on EC2 instances for development, staging and production environments using **k3d** or **kOps**. Developed utility **bash** scripts with **helm** and **kubectl** to manage Kubernetes clusters more efficiently, including configuring **Nginx ingress** controller, cert manager with **Let's encrypt** issuer, **Cillium** Container Network Interface(CNI), IAM Roles for Service Account(IRSA), **Cluster Autoscaler**, AWS Elastic Load Balancer(ELB) or deploying applications like Percona XtraDB Clusters, Keycloak, JupyterHub, Flask and ReactJS based web application, etc.
- * **Ephemeral Worker Clusters:** Designed and developed a worker lifecycle manager using Python within one month to fulfill an **urgent** business requirement. This development **polls** jobs from a MySQL database, then provisions and configures ephemeral EC2 instances by **Packer(pre-build AMI)**, **Terraform** and **cloud-init** to compute jobs **at scale**; implemented AWS S3 mount to significantly reduce raw data downloading **overhead** and added EFS as a file cache for intermediate steps to improve computation **failover**; configured **NVIDIA CUDA toolkit** and **NVIDIA container runtime** for **GPU** workers.
- * **Platform Automation:** To provision or terminate AWS resources using **boto3** or **Terraform**; manage customers' **RBAC** permissions using Keycloak and Github REST API; generating usage and billing report with **AWS S3 Inventory** report, **AWS CloudTrail** and **AWS Cost and Usage** report, made a **Plotly Dash** to analyze cost and usage efficiency.
- * **Jupyterhub:** Configured and maintained Jupyterhub deployment on a Kubernetes cluster with **Node Affinity** to assign pods onto different nodes by requirements and **Cluster Autoscaler** along with **AWS Auto Scaling Group** to accommodate **100+** active users; improved base images' **build time** and maintenance **overhead**.
- * **Observability:** Implemented a small part of the metrics and alerts using **AWS CloudWatch**, and then later integrated **Datadog** for Kubernetes clusters' and ephemeral EC2 instances' metrics and logging through **OpenTelemetry** protocol, synthetic API testing, and UI/UX monitoring.
- * **Security:** Set up codebase **vulnerability** scan with FOSSA; Set up **AWS Secrets Manager** working with **External Secret Store Operator** to secure Kubernetes secrets; Deployed and administrated self-hosted **Keycloak** for **RABC** authentication, further integrated it with **AWS IAM** as an **identity provider** to access AWS resources through **STS**, enabled OpenID Connect(OIDC) authentication flows such as authorization code flow, client credential flow, password grant flow etc.
- * **MySQL Database:** Maintained a self-hosted **Percona XtraDB Clusters** on database **daily backup** stored on **S3**, **mysqldump** backup redundancy, Point-in-Time Recovery(PITR), **deadlock** detection, and slow query log.

Software Engineer(MLOps)

May 2019 - July 2021

dataVediK - Optimize Oil & Gas operations by Machine Learning. [[DrillVedik](#)] [[Details](#)]

Houston, TX

Education

Southern Methodist University, *Master's of Computer Science*

Dallas, TX | Aug 2017 - May 2019

Qingdao University, *Bachelor's of Software Engineering*

Qingdao, China | Aug 2013 - May 2017